

CSCI3230 Fundamentals of Artificial Intelligence

Neural Network Project

Due date: 23:59:59 (GMT +08:00), 8 December, 2011

1. Project Specification

Handwriting recognition is nowadays pre-installed in your iPhone 4S, Galaxy S2, iPad 2 and so on but have you ever thought of the principle and the mechanism behind? In this project, you have a chance to design a letter recognizer based on neural networks to identify 26 capital letters in English from pre-processed handwritten input data. You have to design your neural network, implement a program to train the neural network and output your neural network in a specified format to accomplish the task.

The objective of this project is to identify each of a large number of black-and-white rectangular pixel displays as one of the 26 letters in the English alphabet. The character images were based on 20 different fonts and each letter within these 20 fonts was randomly distorted to produce a file of 20,000 unique stimuli. Each stimulus was converted into 16 numerical attributes that were in turn submitted to our classifier system. A table below summarizes the attributes.

The neural network should learn its weights by using the back-propagation algorithm. You may choose one of the available programming languages, which are listed in the Appendix. However, you are not allowed to use any data mining or machine learning packages.

You are required to output your neural network in a format as described below and name the output file *best.nn* for submission; Also, you have to submit a *preprocessor* program to convert the raw data to the inputs to neural network; a *trainer* program to train the neural network (*best.nn*) and *two shell scripts* to compile and run the preprocessor and the trainer program respectively.

2. Dataset

You are provided with the *raw-training.name* and *raw-training.data* (same dataset format with Weka Laboratory) which you can find in the course website. The raw data file contains a number of records. Each record represents a letter and has seventeen items.

	Attribute	Type	Description
1	<i>letter</i>	Character	Capital letter (26 values from A to Z)
2	<i>x-box</i>	Integer	Horizontal position of box
3	<i>y-box</i>	Integer	Vertical position of box
4	<i>width</i>	Integer	Width of box
5	<i>high</i>	Integer	Height of box
6	<i>onpix</i>	Integer	Total # on pixels
7	<i>x-bar</i>	Integer	Mean x of on pixels in box
8	<i>y-bar</i>	Integer	Mean y of on pixels in box
9	<i>x2bar</i>	Integer	Mean x variance
10	<i>y2bar</i>	Integer	Mean y variance
11	<i>xybar</i>	Integer	Mean x y correlation
12	<i>x2ybr</i>	Integer	Mean of $x * x * y$
13	<i>xy2br</i>	Integer	Mean of $x * y * y$
14	<i>x-ege</i>	Integer	Mean edge count left to right
15	<i>xegvy</i>	Integer	Correlation of x-edge with y
16	<i>y-ege</i>	Integer	Mean edge count bottom to top
17	<i>yegvx</i>	Integer	Correlation of y-edge with x

3. The Neural Network and File Format of *best.nn*

As you are suggested to do data preprocessing, we do NOT limit the number of input neurons, the number of hidden layers and the number of neurons in each hidden layer. You can choose them as appropriate.

However, we have set the number of output neurons to be **26**. 'A', 'B', ..., 'Z' are represented from left to right. For each output neuron, if its output is > 0.5 , it is considered to be 1, otherwise 0. The output of your neural network should be in 1-of-C encoding, which means one and only one output neuron should be > 0.5 for each record, and the digit represented by that neuron is considered to be the classified class. For example, if your neural network outputs [0.1, 0.6, 0.1, 0.2, 0.3, 0.4, 0.5, 0.1, 0.2, 0.3, 0.4, 0.5, 0.1, 0.2, 0.3, 0.4, 0.5, 0.1, 0.2, 0.3, 0.4, 0.5, 0.1, 0.2, 0.3, 0.4], it is considered to be [0 1 0, ..., 0], and your neural network classifies the record to be an 'B'. Note that if more than one or no neuron output is > 0.5 , that record is considered to be **misclassified**.

The activation function used for a neuron is the sigmoid function, which is

$$g(t) = \frac{1}{1 + e^{-t}}$$

The possible structure of neural network allowed is illustrated in figure 1.




Figure 1

The first line of the output should be the classification accuracy (0 to 1.0), the second line should be the number of hidden layers, ie. k , and the rest store the weights, divided into $k+1$ sections, as follows:

Each section stores the weights and bias terms between two successive layers i and $i+1$ (Input to layer 1, layer 1 to layer 2, ..., layer k to Output). Each section starts with a line with 4 items, separated with space:

$$label_i \ H_i \ label_{i+1} \ H_{i+1}$$

where $label_i$ is the label of layer i ('I' for input layer, 'H' for hidden layer, 'O' for output layer), H_i is the number of neurons in layer i , $label_{i+1}$ is the label of layer $i+1$, H_{i+1} is the number of neurons in layer $i+1$. Then each of the next H_{i+1} lines should store the bias term and weights of the r -th neuron in layer $i+1$. Each line should look like:

$$w_{i,0} \ w_{i,1} \ w_{i,2} \ \dots \ w_{i,H_i}$$

where $w_{i,0}$ is the bias term, (which acts like a weight of an input of constant -1) the following H_i numbers $w_{i,j}$ should be the weight from the j -th neuron in layer i to the r -th neuron in layer $i+1$.




Figure 2

For a complete example, for the neural network in figure 2 with classification accuracy of 0.9, the output file would look like:

```
0.9
2
I 3 H 2
0.1 0.3 0.4 0.5
0.2 -0.6 0.7 -0.8
H 2 H 3
0.2 0.1 -0.2
0.1 0.5 0.6
0.3 0.3 0.4
H 3 O 2
0.7 0.7 0.8 0.9
0.8 0.5 0.4 0.3
```

4. Preprocessing Program

In this project, you are given only raw data. You are advised to do preprocessing on the raw data to help building a good neural network. But you **MUST** submit a preprocessor program which converts raw data without class labels (**class labels would appear as '?' when missing**) into inputs to the neural network. (With the same relative order ie. do not shuffle the data records) The format is: for raw data with N records to be classified, there should be $N+1$ line(s) in the output. The first line is the number of records and the number of features, i.e. N and H_0 . Each of the remaining lines corresponds to a record and should have H_0 numbers, to be used as the H_0 inputs to the neural network, where H_0 is the number of input neurons you have chosen. For example, if the raw data contains 2 records, and you choose to have 3 inputs neurons, the file may look like:

```
2 3
5 1 8
6 7 3
```

In order for your trainer program to train the neural network, an answer file containing only class labels has to be generated by your preprocessor program. The first line is the number of records, i.e. N . Each of the remaining lines is the class label for the corresponding record. For the above example, the answer file may look like:

```
2
EI
N
```

5. Usage and Format of Shell Scripts

In this project, you are required to submit 2 shell scripts to compile and run the preprocessor and trainer programs, named "*preprocessor.sh*" and "*trainer.sh*" respectively.

For *preprocessor.sh*:

The required execution command is:

```
./preprocessor.sh raw-training.data dataset.dat answer.dat
```

where *raw-training.dat* is the raw dataset that you can find in our course homepage;

dataset.dat is the preprocessed data, the inputs of the neural network;

answer.dat is the class label for the preprocessed data

Example (if your preprocessor is written in C language):

```
#!/bin/sh
gcc -o preprocess preprocess.c -Wall -lm
./preprocess $1 $2 $3
```

For *trainer.sh*:

The required execution command is:

```
./trainer.sh dataset.dat answer.dat best.nn <max_iteration> <small_sample_mode>
```

where *dataset.dat* is the preprocessed data generated by *preprocessor.sh* as input;

answer.dat is the class label for the preprocessed data;

best.nn is the neural network your program trained as output;

<max_iteration> is the maximum number of iterations for the learning process[†];

<small_sample_mode> is the flag indicating the mode of training[†]

Example (if your trainer is written in Java language):

```
#!/bin/sh
javac trainer.java
java trainer $1 $2 $4 $5 > $3
```

[†] In order for the tutors to examine the correctness and performance of your backward propagation algorithm, two extra arguments are required to be accepted by your trainer program. The first one is `<max_iteration>`, the maximum number of iteration for the learning process, it is one of the early stopping criteria. The second one is `<small_sample_mode>`, the flag indicates the training mode, 1 for small sample mode and 0 for typical mode. During small sample mode, your trainer program should be able to handle a tiny dataset with 260 observations only.

6. Grading Criteria

We will test your neural network on the training data set and a secret testing data set. We will use your preprocessing program to get the inputs to be fed to your neural network, then we will simulate the neural network using our own program to calculate the classification accuracy.

The division of marks is:

- | | |
|---|-----|
| - Implementation of back-propagation | 50% |
| - Classification accuracy of Training dataset | 20% |
| - Classification accuracy of Testing dataset | 30% |

7. Assignment Submission

Zip the following with filename as “NN.zip”:

- Source code of your preprocessor program
- Source code of your trainer program
- Your best neural network named “*best.nn*”
- A shell script to compile and run the preprocessor program
- A shell script to compile and run the trainer program
- A readme file that explain how to use your shell scripts and programs

You **MUST** submit the ZIP file to the submission system on our course homepage (within CUHK network), otherwise, we will **NOT** mark your assignment.

8. Important Points

You **MUST STRICTLY** follow these points:

- You **MUST** strictly follow the submission guidelines.
- Your trainer is only given 1 MINUTE of execution time.
- Late submission will **NOT** be entertained according to our submission system settings.
- Plagiarism will be **SERIOUSLY** punished.

9. Late Submission

According to the course homepage, late submission will lead to marks deduction.

No. of Days Late	Marks Deduction
1	10%
2	30%
3	60%
4 or above	100%

Appendix

Language	Version	Usage
C	gcc version 4.3.3 (Ubuntu 4.3.3-5ubuntu4)	<p>Usage: gcc [options] file...</p> <p>Options:</p> <ul style="list-style-type: none"> -pass-exit-codes Exit with highest error code from a phase --help Display this information --target-help Display target specific command line options --help={target optimizers warnings undocumented params}[,{[[^]]joined [[^]]separate}] Display specific types of command line options <p>(Use '-v --help' to display command line options of sub-processes)</p> <ul style="list-style-type: none"> -dumpspecs Display all of the built in spec strings -dumpversion Display the version of the compiler -dumpmachine Display the compiler's target processor -print-search-dirs Display the directories in the compiler's search path -print-libgcc-file-name Display the name of the compiler's companion library -print-file-name=<lib> Display the full path to library <lib> -print-prog-name=<prog> Display the full path to compiler component <prog> -print-multi-directory Display the root directory for versions of libgcc -print-multi-lib Display the mapping between command line options and multiple library search directories -print-multi-os-directory Display the relative path to OS libraries -print-sysroot-headers-suffix Display the sysroot suffix used to find headers -Wa,<options> Pass comma-separated <options> on to the assembler -Wp,<options> Pass comma-separated <options> on to the preprocessor -Wl,<options> Pass comma-separated <options> on to the linker -Xassembler <arg> Pass <arg> on to the assembler -Xpreprocessor <arg> Pass <arg> on to the preprocessor -Xlinker <arg> Pass <arg> on to the linker -combine Pass multiple source files to compiler at once -save-temps Do not delete intermediate files -pipe Use pipes rather than intermediate files -time Time the execution of each subprocess -specs=<file> Override built-in specs with the contents of <file> -std=<standard> Assume that the input sources are for <standard> --sysroot=<directory> Use <directory> as the root directory for headers and libraries -B <directory> Add <directory> to the compiler's search paths -b <machine> Run gcc for target <machine>, if installed -V <version> Run gcc version number <version>, if installed -v Display the programs invoked by the compiler -### Like -v but options quoted and commands not executed -E Preprocess only; do not compile, assemble or link -S Compile only; do not assemble or link -c Compile and assemble, but do not link

		<p>-o <file> Place the output into <file></p> <p>-x <language> Specify the language of the following input files Permissible languages include: c c++ assembler none 'none' means revert to the default behavior of guessing the language based on the file's extension</p> <p>Options starting with -g, -f, -m, -O, -W, or --param are automatically passed on to the various sub-processes invoked by gcc. In order to pass other options on to these processes the -W<letter> options must be used.</p> <p>For bug reporting instructions, please see: <file:///usr/share/doc/gcc-4.3/README.Bugs>.</p>
C++	gcc version 4.3.3 (Ubuntu 4.3.3-5ubuntu4)	<p>Usage: g++ [options] file...</p> <p>Options:</p> <ul style="list-style-type: none"> -pass-exit-codes Exit with highest error code from a phase --help Display this information --target-help Display target specific command line options --help={target optimizers warnings undocumented params}[,{[[^]]joined [[^]]separate}] Display specific types of command line options (Use '-v --help' to display command line options of sub-processes) -dumpspecs Display all of the built in spec strings -dumpversion Display the version of the compiler -dumpmachine Display the compiler's target processor -print-search-dirs Display the directories in the compiler's search path -print-libgcc-file-name Display the name of the compiler's companion library -print-file-name=<lib> Display the full path to library <lib> -print-prog-name=<prog> Display the full path to compiler component <prog> -print-multi-directory Display the root directory for versions of libgcc -print-multi-lib Display the mapping between command line options and multiple library search directories -print-multi-os-directory Display the relative path to OS libraries -print-sysroot-headers-suffix Display the sysroot suffix used to find headers -Wa,<options> Pass comma-separated <options> on to the assembler -Wp,<options> Pass comma-separated <options> on to the preprocessor -Wl,<options> Pass comma-separated <options> on to the linker -Xassembler <arg> Pass <arg> on to the assembler -Xpreprocessor <arg> Pass <arg> on to the preprocessor -Xlinker <arg> Pass <arg> on to the linker -combine Pass multiple source files to compiler at once -save-temps Do not delete intermediate files -pipe Use pipes rather than intermediate files -time Time the execution of each subprocess -specs=<file> Override built-in specs with the contents of <file> -std=<standard> Assume that the input sources are for <standard> --sysroot=<directory> Use <directory> as the root directory for headers and libraries

		-help Print a synopsis of standard options -Akey[=value] Options to pass to annotation processors -X Print a synopsis of nonstandard options -J<flag> Pass <flag> directly to the runtime system
swi-Prolog	SWI-Prolog (Multi-threaded, 32 bits, Version 5.6.59)	swipl: Usage: 1) swipl --help Display this message (also -h) 2) swipl --version Display version information (also -v) 3) swipl --arch Display architecture 4) swipl --dump-runtime-variables[=format] Dump link info in sh(1) format 5) swipl [options] 6) swipl [options] [-o output] -c file ... 7) swipl [options] [-o output] -b bootfile -c file ... Options: -x state Start from state (must be first) -[LGTA]size[KMG] Specify {Local,Global,Trail,Argument} limits -t toplevel Toplevel goal -g goal Initialisation goal -f file User initialisation file -F file System initialisation file -s file Script source file [+/-]tty Allow tty control -O Optimised compilation --nosignals Do not modify any signal handling --nodebug Omit generation of debug info --quiet Quiet operation (also -q) --home=DIR Use DIR as SWI-Prolog home
CLisp	GNU CLISP 2.44.1 (2008-02-23) (built 3436675404)	GNU CLISP (http://clisp.cons.org/) is an ANSI Common Lisp. Usage: clisp [options] [lispfile [argument ...]] When 'lispfile' is given, it is loaded and '*ARGS*' is set to the list of argument strings. Otherwise, an interactive read-eval-print loop is entered. Informative output: -h, --help - print this help and exit --version - print the version information --license - print the licensing information -help-image - print image-specific help and exit Memory image selection: -B lisplibdir - set the installation directory -K linkingset - use this executable and memory image -M memfile - use this memory image -m size - memory size (size = xxxxxxxB or xxxKB or xMB) Internationalization: -L language - set user language -N nlsdir - NLS catalog directory -Edomain encoding - set encoding

		<p>Interoperability:</p> <ul style="list-style-type: none"> -q, --quiet, --silent, -v, --verbose - verbosity level: affects banner, *LOAD-VERBOSE*/*COMPILE-VERBOSE*, and *LOAD-PRINT*/*COMPILE-PRINT* -w - wait for a keypress after program termination -I - be ILISP-friendly <p>Startup actions:</p> <ul style="list-style-type: none"> -ansi - more ANSI CL compliance -traditional - traditional (undoes -ansi) -modern - start in a case-sensitive lowercase-preferring package -p package - start in the package -C - set *LOAD-COMPILING* to T -norc - do not load the user ~/.clisprc file -i file - load initfile (can be repeated) <p>Actions:</p> <ul style="list-style-type: none"> -c [-l] lispfile [-o outputfile] - compile lispfile -x expressions - execute the expressions, then exit <p>Depending on the image, positional arguments can mean:</p> <ul style="list-style-type: none"> lispfile [argument ...] - load script, then exit [argument ...] - run the init-function <p>arguments are placed in EXT:*ARGS* as strings.</p> <p>These actions put CLISP into a batch mode, which is overridden by</p> <ul style="list-style-type: none"> -on-error action - action can be one of debug, exit, abort, appease -repl - enter the interactive read-eval-print loop when done <p>Default action is an interactive read-eval-print loop.</p>
Python	<p>Python 2.6.2 (release26-maint, Apr 19 2009, 01:56:41)</p>	<p>usage: python [option] ... [-c cmd -m mod file -] [arg] ...</p> <p>Options and arguments (and corresponding environment variables):</p> <ul style="list-style-type: none"> -B : don't write .py[co] files on import; also PYTHONDONTWRITEBYTECODE=x -c cmd : program passed in as string (terminates option list) -d : debug output from parser; also PYTHONDEBUG=x -E : ignore PYTHON* environment variables (such as PYTHONPATH) -h : print this help message and exit (also --help) -i : inspect interactively after running script; forces a prompt even if stdin does not appear to be a terminal; also PYTHONINSPECT=x -m mod : run library module as a script (terminates option list) -O : optimize generated bytecode slightly; also PYTHONOPTIMIZE=x -OO : remove doc-strings in addition to the -O optimizations -Q arg : division options: -Qold (default), -Qwarn, -Qwarnall, -Qnew -s : don't add user site directory to sys.path; also PYTHONNOUSERSITE -S : don't imply 'import site' on initialization -t : issue warnings about inconsistent tab usage (-tt: issue errors) -u : unbuffered binary stdout and stderr; also PYTHONUNBUFFERED=x see man page for details on internal buffering relating to '-u' -v : verbose (trace import statements); also PYTHONVERBOSE=x can be supplied multiple times to increase verbosity -V : print the Python version number and exit (also --version)

		<p>-W arg : warning control; arg is action:message:category:module:lineno</p> <p>-x : skip first line of source, allowing use of non-Unix forms of #!cmd</p> <p>-3 : warn about Python 3.x incompatibilities that 2to3 cannot trivially fix</p> <p>file : program read from script file</p> <p>- : program read from stdin (default; interactive mode if a tty)</p> <p>arg ...: arguments passed to program in sys.argv[1:]</p> <p>Other environment variables:</p> <p>PYTHONSTARTUP: file executed on interactive startup (no default)</p> <p>PYTHONPATH : ':'-separated list of directories prefixed to the default module search path. The result is sys.path.</p> <p>PYTHONHOME : alternate <prefix> directory (or <prefix>:<exec_prefix>). The default module search path uses <prefix>/pythonX.X.</p> <p>PYTHONCASEOK : ignore case in 'import' statements (Windows).</p> <p>PYTHONIOENCODING: Encoding[:errors] used for stdin/stdout/stderr.</p>
Ruby	ruby 1.8.7 (2008-08-11 patchlevel 72) [i486-linux]	<p>Usage: ruby [switches] [--] [programfile] [arguments]</p> <p>-0[octal] specify record separator (\0, if no argument)</p> <p>-a autosplit mode with -n or -p (splits \$_ into \$F)</p> <p>-c check syntax only</p> <p>-Cdirectory cd to directory, before executing your script</p> <p>-d set debugging flags (set \$DEBUG to true)</p> <p>-e 'command' one line of script. Several -e's allowed. Omit [programfile]</p> <p>-Fpattern split() pattern for autosplit (-a)</p> <p>-i[extension] edit ARGV files in place (make backup if extension supplied)</p> <p>-Idirectory specify \$LOAD_PATH directory (may be used more than once)</p> <p>-Kkcode specifies KANJI (Japanese) code-set</p> <p>-l enable line ending processing</p> <p>-n assume 'while gets(); ... end' loop around your script</p> <p>-p assume loop like -n but print line also like sed</p> <p>-rlibrary require the library, before executing your script</p> <p>-s enable some switch parsing for switches after script name</p> <p>-S look for the script using PATH environment variable</p> <p>-T[level] turn on tainting checks</p> <p>-v print version number, then turn on verbose mode</p> <p>-w turn warnings on for your script</p> <p>-W[level] set warning level; 0=silence, 1=medium, 2=verbose (default)</p> <p>-x[directory] strip off text before #!ruby line and perhaps cd to directory</p> <p>--copyright print the copyright</p> <p>--version print the version</p>
Perl	perl, v5.10.0 built for i486-linux-gnu-thread-multi	<p>Usage: perl [switches] [--] [programfile] [arguments]</p> <p>-0[octal] specify record separator (\0, if no argument)</p> <p>-a autosplit mode with -n or -p (splits \$_ into @F)</p> <p>-C[number/list] enables the listed Unicode features</p> <p>-c check syntax only (runs BEGIN and CHECK blocks)</p> <p>-d[:debugger] run program under debugger</p> <p>-D[number/list] set debugging flags (argument is a bit mask or alphabets)</p>

		-e program -E program -f -F/pattern/ -i[extension] -Idirectory -l[octal] -[mM] [-]module -n -p -P -s -S -t -T -u -U -v -V[:variable] -w -W -x[directory] -X	one line of program (several -e's allowed, omit programfile) like -e, but enables all optional features don't do \$sitelib/sitecustomize.pl at startup split() pattern for -a switch (//'s are optional) edit <> files in place (makes backup if extension supplied) specify @INC/#include directory (several -I's allowed) enable line ending processing, specifies line terminator execute "use/no module..." before executing program assume "while (<>) { ... }" loop around program assume loop like -n but print line also, like sed run program through C preprocessor before compilation enable rudimentary parsing for switches after programfile look for programfile using PATH environment variable enable tainting warnings enable tainting checks dump core after parsing program allow unsafe operations print version, subversion (includes VERY IMPORTANT perl info) print configuration summary (or a single Config.pm variable) enable many useful warnings (RECOMMENDED) enable all warnings strip off text before #!perl line and perhaps cd to directory disable all warnings
--	--	---	---